

# MA 578 HW5 Solutions

We will be using the following packages and definitions:

```
library(rstan)
library(bayesplot)
library(beanplot)
library(tidyverse)

# sample from inverse (scaled) chi-square with parameters `nu` and `tau2`;
# nu_tau2 = nu * tau2 for convenience
rinvchisq <- function (ns, nu, nu_tau2) 1 / rgamma(ns, nu / 2, nu_tau2 / 2)

plot_hist <- function (x, ...)
  hist(x, col = "gray", border = "white", main = "", ...)

mcmc_array <- function (ns, nchains = 1, params) {
  nparams <- length(params)
  array(dim = c(ns, nchains, nparams),
        dimnames = list(iterations = NULL,
                          chains = paste0("chain", 1:nchains),
                          parameters = params))
}

iter <- 2000
warmup <- floor(iter / 2)
nsamples <- iter - warmup
nchains <- 4
```

## 1 (BDA 11.2)

Using random walk Metropolis-Hastings with non-informative prior  $\mathbb{P}(\alpha, \beta) \propto 1$ .

```
bioassay <- read.csv("data/bioassay.csv", comment = "#")
n <- bioassay$n; y <- bioassay$deaths; x <- bioassay$logdose
log_p <- function (alpha, beta) { # log posterior
  p <- plogis(alpha + beta * x)
  sum(y * log(p) + (n - y) * log(1 - p))
}

sigma_alpha <- .25
sigma_beta <- 5 * sigma_alpha
nthin <- 10
sims <- mcmc_array(nsamples, nchains, c("alpha", "beta", "lp_"))

gb <- glm(cbind(y, n - y) ~ x, family = binomial)
alpha <- coef(gb)[1]; beta <- coef(gb)[2]
for (chain in 1:nchains) {
  for (it in 1:iter) {
    for (thin in 1:nthin) {
      alpha_c <- rnorm(1, alpha, sigma_alpha)
      log_r <- log_p(alpha_c, beta) - log_p(alpha, beta)
      if (log_r >= 0 || log(runif(1)) <= log_r) alpha <- alpha_c # accept?
    }
  }
}
```

```

    beta_c <- rnorm(1, beta, sigma_beta)
    log_r <- log_p(alpha, beta_c) - log_p(alpha, beta)
    if (log_r >= 0 || log(runif(1)) <= log_r) beta <- beta_c # accept?
  }

  if (it > warmup)
    sims[it - warmup, chain, ] <- c(alpha, beta, log_p(alpha, beta))
}
}

monitor(sims, warmup = 0)

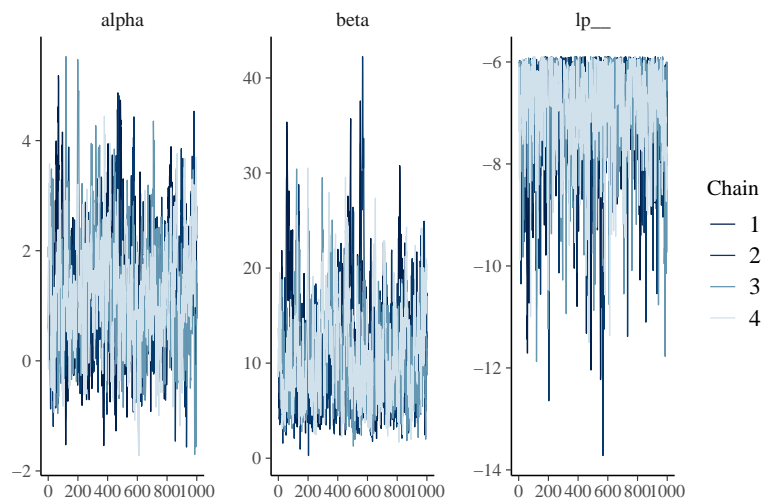
```

Inference for the input samples (4 chains: each with iter = 1000; warmup = 0):

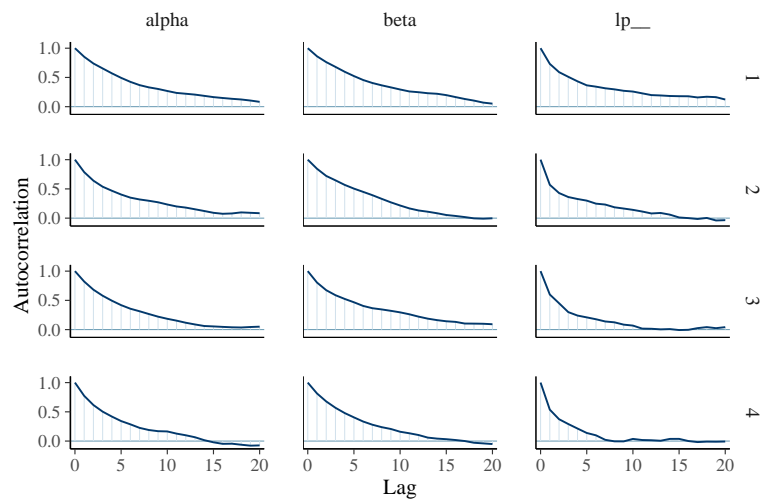
	Q5	Q50	Q95	Mean	SD	Rhat	Bulk_ESS	Tail_ESS
alpha	-0.4	1.2	3.1	1.2	1.1	1.01	349	484
beta	4.2	10.1	21.4	11.1	5.4	1.01	350	473
lp__	-8.9	-6.6	-5.9	-6.9	1.0	1.00	604	582

For each parameter, Bulk\_ESS and Tail\_ESS are crude measures of effective sample size for bulk and tail quantities respectively (an ESS > 100 per chain is considered good), and Rhat is the potential scale reduction factor on rank normalized split chains (at convergence, Rhat <= 1.05).

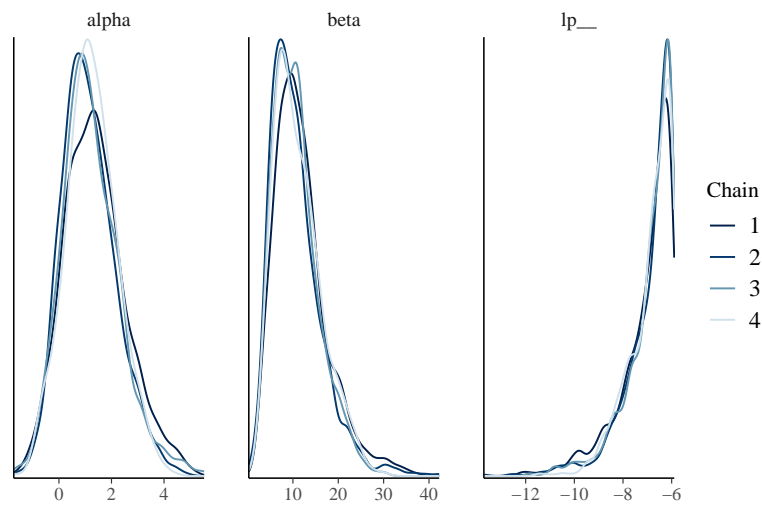
```
mcmc_trace(sims)
```



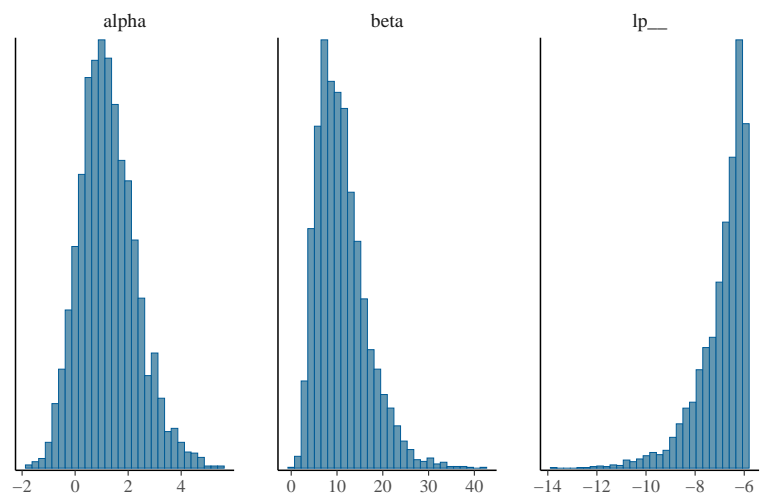
```
mcmc_acf(sims)
```



```
mcmc_dens_overlay(sims)
```



```
mcmc_hist(sims)
```



For comparison, here's the equivalent Stan output:

```

N <- length(y)
sm <- stan_model(model_code = "data { int<lower=0> N; int<lower=0> n[N];
  int<lower=0> y[N]; vector[N] x; }
parameters { real alpha; real beta; }
model { y ~ binomial_logit(n, alpha + beta * x); }")
sf <- sampling(sm, data = c("N", "n", "y", "x"))

```

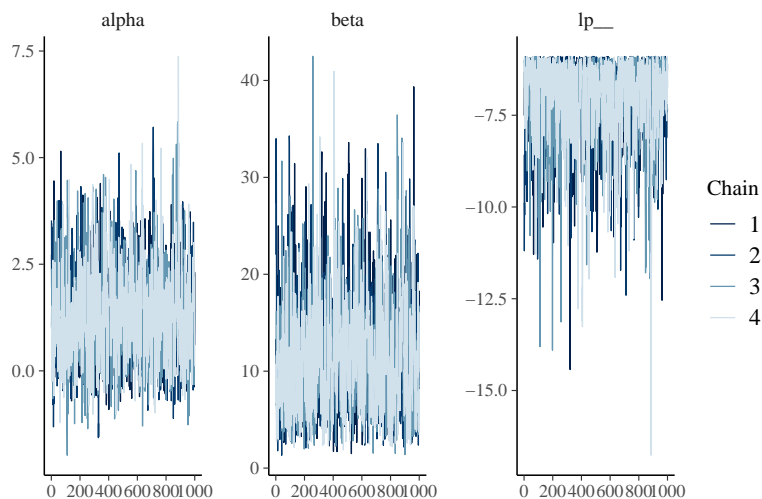
```
monitor(sf)
```

Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):

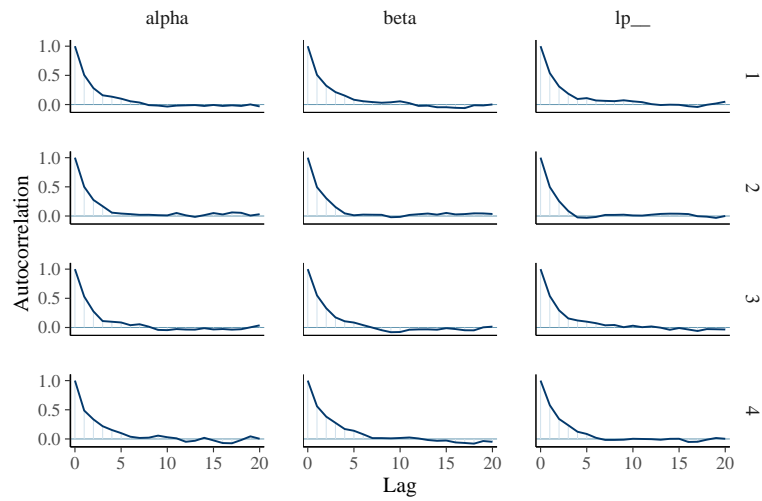
	Q5	Q50	Q95	Mean	SD	Rhat	Bulk_ESS	Tail_ESS
alpha	-0.3	1.2	3.2	1.3	1.1	1	1191	1496
beta	4.2	10.8	22.3	11.6	5.7	1	1181	1281
lp__	-9.2	-6.7	-5.9	-7.0	1.1	1	1281	1501

For each parameter, Bulk\_ESS and Tail\_ESS are crude measures of effective sample size for bulk and tail quantities respectively (an ESS > 100 per chain is considered good), and Rhat is the potential scale reduction factor on rank normalized split chains (at convergence, Rhat <= 1.05).

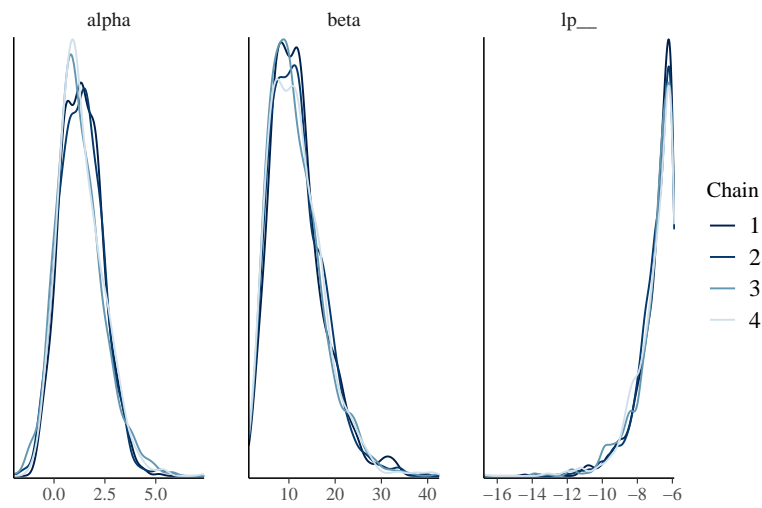
```
mcmc_trace(sf)
```



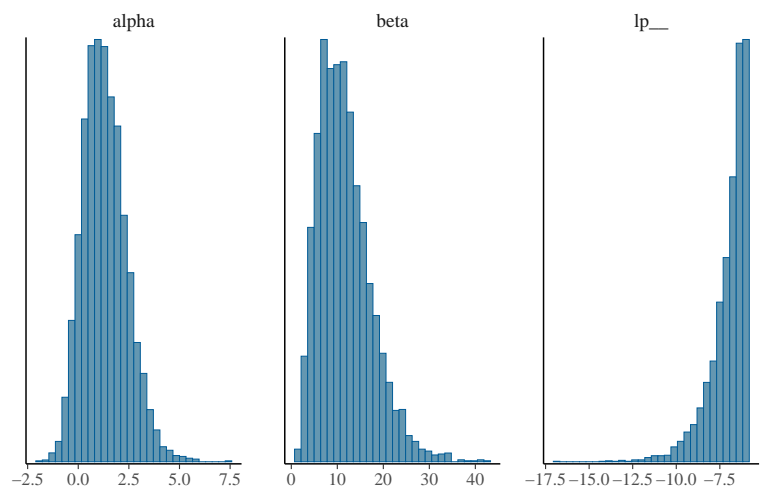
```
mcmc_acf(sf)
```



```
mcmc_dens_overlay(sf)
```



```
mcmc_hist(sf)
```



## 2 (BDA 11.3)

We have three analyses—separate, pooled, and hierarchical (partially pooled)—of a set of quality control measurements. Here  $y_{ij}$  is the  $i$ -th measurement, out of  $n_j$  measurements, in the  $j$ -th machine from a total of  $J = 6$  machines.

```
machine <- read.csv("data/machine.csv", comment = "#")
machine %>% group_by(machine) %>% summarize(qc = paste(qc, collapse = ", "))
```

```
# A tibble: 6 x 2
```

```
  machine qc
  <int> <chr>
1     1 83, 92, 92, 46, 67
2     2 117, 109, 114, 104, 87
3     3 101, 93, 92, 86, 67
4     4 105, 119, 116, 102, 116
5     5 79, 97, 103, 79, 92
6     6 57, 92, 104, 77, 100
```

```
# sufficient statistics:
```

```
nj <- with(machine, tapply(qc, machine, length))
yj <- with(machine, tapply(qc, machine, mean))
s2j <- with(machine, tapply(qc, machine, var))
S <- sum(s2j * (nj - 1))
J <- length(yj)
n <- sum(nj)
```

For each analysis, we want (i)  $\theta_6 | y$ , (ii)  $\tilde{y}_6 | y$ , and (iii)  $\theta_7 | y$ . These posterior distributions should be estimated using Gibbs sampling, as requested.

### Separate

The model we assume is  $y_{ij} | \theta_j, \sigma^2 \stackrel{\text{ind}}{\sim} N(\theta_j, \sigma^2)$  for  $i = 1, \dots, n_j$ ,  $j = 1, \dots, J$ , and non-informative priors  $\mathbb{P}(\theta, \log \sigma) \propto 1$ . For the Gibbs sampler, since  $\bar{y}_j | \theta_j, \sigma^2 \stackrel{\text{ind}}{\sim} N(\theta_j, \sigma^2/n_j)$ , we can first see that  $\theta_j | \sigma^2, y \stackrel{\text{ind}}{\sim} N(\bar{y}_j, \sigma^2/n_j)$ . The joint posterior is

$$\mathbb{P}(\theta, \sigma^2 | y) \propto (\sigma^2)^{-\sum_j n_j/2} \exp \left\{ -\frac{\sum_{i,j} (y_{ij} - \theta_j)^2}{2\sigma^2} \right\} (\sigma^2)^{-1}$$

and so  $\sigma^2 | \theta, y \sim \text{Inv-}\chi^2(\sum_j n_j, \sum_{i,j} (y_{ij} - \theta_j)^2 / \sum_j n_j)$ .

```
sims <- mcmc_array(nsamples, nchains, c(paste0("theta", 1:J), "sigma", "lp__"))
```

```
sigma2 <- var(machine$qc)
for (chain in 1:nchains) {
  for (it in 1:iter) {
    theta <- rnorm(J, yj, sqrt(sigma2 / nj))
    nt <- S + sum(nj * (yj - theta) ^ 2)
    sigma2 <- rinvcchisq(1, n, nt)
    target <- -(n / 2 + 1) * log(sigma2) - nt / (2 * sigma2)
    if (it > warmup)
      sims[it - warmup, chain, ] <- c(theta, sqrt(sigma2), target)
  }
}

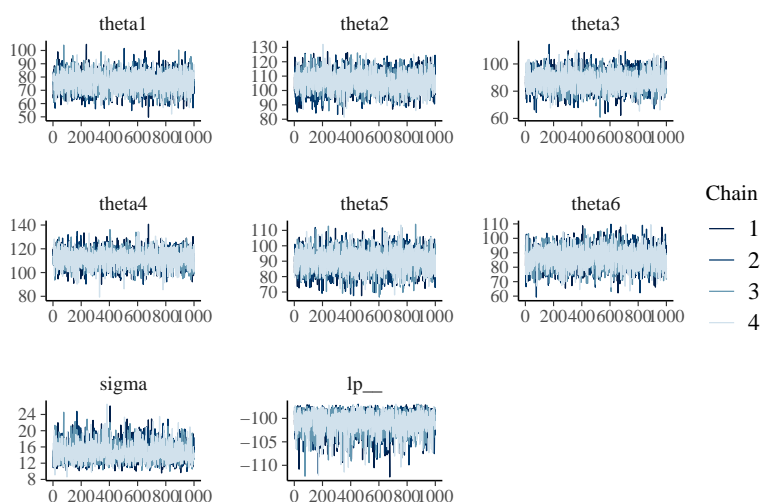
monitor(sims, warmup = 0)
```

Inference for the input samples (4 chains: each with iter = 1000; warmup = 0):

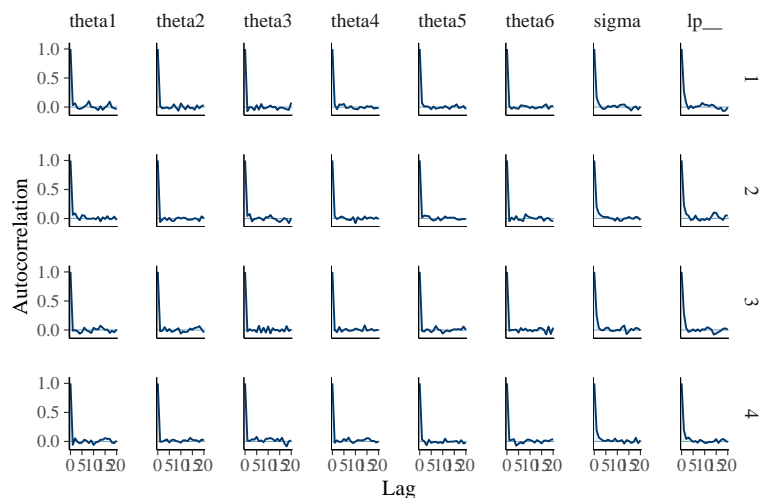
	Q5	Q50	Q95	Mean	SD	Rhat	Bulk_ESS	Tail_ESS
theta1	65.0	76.1	87.0	76.1	6.7	1	3628	3658
theta2	95.3	106.2	116.9	106.2	6.6	1	4146	3727
theta3	77.1	87.8	98.5	87.9	6.6	1	3912	3674
theta4	100.8	111.6	122.1	111.5	6.6	1	3712	3560
theta5	79.2	89.8	100.8	90.0	6.6	1	3677	3631
theta6	75.2	85.9	97.0	86.1	6.7	1	4023	3924
sigma	11.6	14.5	18.8	14.8	2.2	1	2500	2995
lp__	-104.9	-100.4	-98.0	-100.8	2.2	1	2527	3086

For each parameter, Bulk\_ESS and Tail\_ESS are crude measures of effective sample size for bulk and tail quantities respectively (an ESS > 100 per chain is considered good), and Rhat is the potential scale reduction factor on rank normalized split chains (at convergence, Rhat <= 1.05).

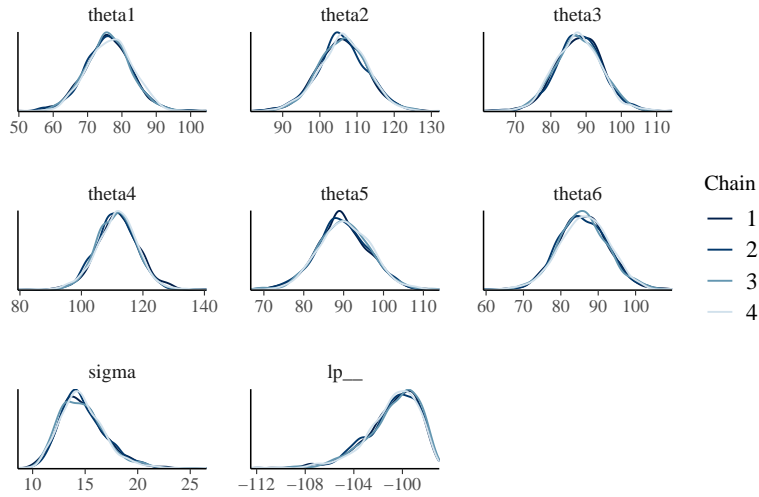
`mcmc_trace(sims)`



`mcmc_acf(sims)`

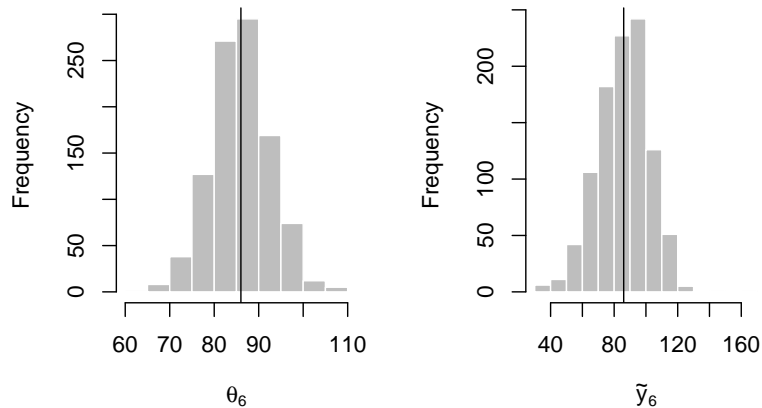


```
mcmc_dens_overlay(sims)
```



Now we can provide what has been asked. Since we are assuming a separate model, the posterior distribution for  $\theta_7$  is the same as its prior, flat, due to no borrowing of information from other machines. Not surprisingly, the posteriors are in good agreement with the data.

```
# we can just use the first chain
theta6 <- sims[, 1, 6]; sigma <- sims[, 1, 7]
ypred6 <- rnorm(nsamples, theta6, sigma)
op <- par(mfrow = c(1, 2))
plot_hist(theta6, xlab = expression(theta[6])); abline(v = yj[6])
plot_hist(ypred6, xlab = expression(tilde(y)[6])); abline(v = yj[6])
```



```
par(op)
```

For comparison, here's the Stan equivalent model; the results are not shown to save space, but they are equivalent to Gibbs sampling.

```
y <- machine$qc; idx <- c(0, cumsum(nj))[-J]
sm <- stan_model(model_code = "data {
  int<lower=0> J; int<lower=0> nj[J]; int<lower=0> idx[J]; int<lower=0> n; real y[n];
}
parameters { vector[J] theta; real<lower=0> sigma; }
model {
  for (j in 1:J) for (i in 1:nj[j]) y[idx[j] + i] ~ normal(theta[j], sigma);
  target += -log(sigma);
}
```



```
})
sf <- sampling(sm, data = c("J", "nj", "idx", "n", "y"))
```

## Pooled

Here we have  $y_{ij} | \mu, \tau^2 \stackrel{\text{iid}}{\sim} N(\mu, \tau^2)$ , with non-informative priors  $\mathbb{P}(\mu, \log \tau) \propto 1$ . The joint posterior is

$$\mathbb{P}(\mu, \tau^2 | y) \propto (\tau^2)^{-\left(\sum_j n_j/2 + 1\right)} \exp \left\{ -\frac{\sum_{i,j} (y_{ij} - \mu)^2}{2\tau^2} \right\}.$$

Thus, we can see that  $\mu | \tau^2, y \sim N(\bar{y}, \tau^2 / \sum_j n_j)$ , where  $\bar{y} = \sum_{i,j} y_{ij} / \sum_j n_j$  is the overall mean, and  $\tau^2 | \mu, y \sim \text{Inv-}\chi^2(\sum_j n_j, \sum_{i,j} (y_{ij} - \mu)^2 / \sum_j n_j)$ .

```
ybar <- weighted.mean(yj, nj); s2 <- (n - 1) * var(machine$qc)
sims <- mcmc_array(nsamples, nchains, c("mu", "tau", "lp_"))

tau2 <- s2 / (n - 1)
for (chain in 1:nchains) {
  for (it in 1:iter) {
    mu <- rnorm(1, ybar, sqrt(tau2 / n))
    nt <- s2 + n * (ybar - mu) ^ 2
    tau2 <- rinvcisq(1, n, nt)
    target <- -(n / 2 + 1) * log(tau2) - nt / (2 * tau2)
    if (it > warmup)
      sims[it - warmup, chain, ] <- c(mu, sqrt(tau2), target)
  }
}

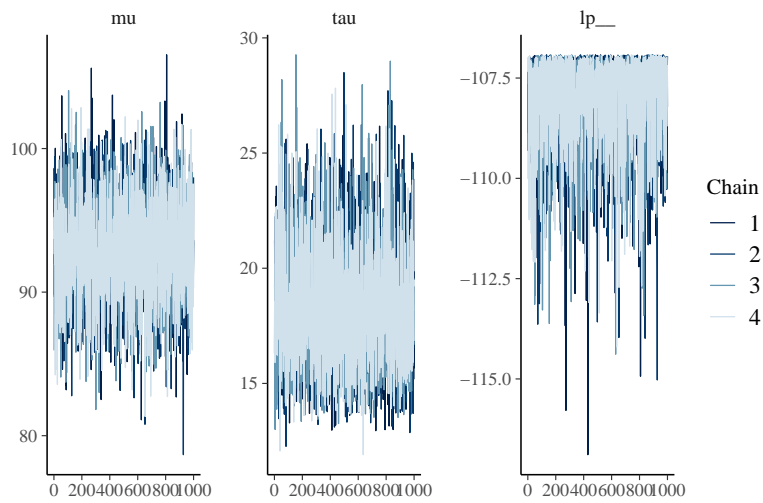
monitor(sims, warmup = 0)
```

Inference for the input samples (4 chains: each with iter = 1000; warmup = 0):

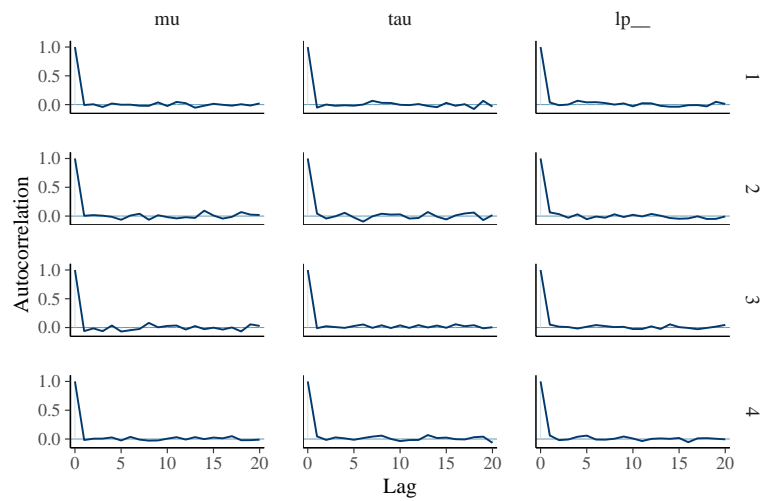
	Q5	Q50	Q95	Mean	SD	Rhat	Bulk_ESS	Tail_ESS
mu	87.4	92.9	98.3	92.9	3.3	1	4178	3695
tau	14.8	18.1	22.9	18.4	2.5	1	3908	3831
lp_	-110.3	-107.7	-107.0	-108.0	1.1	1	3752	3621

For each parameter, Bulk\_ESS and Tail\_ESS are crude measures of effective sample size for bulk and tail quantities respectively (an ESS > 100 per chain is considered good), and Rhat is the potential scale reduction factor on rank normalized split chains (at convergence, Rhat <= 1.05).

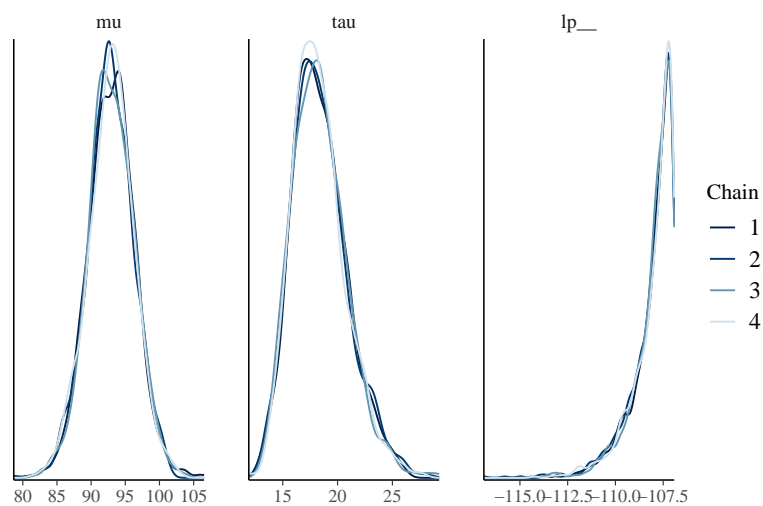
```
mcmc_trace(sims)
```



`mcmc_acf(sims)`

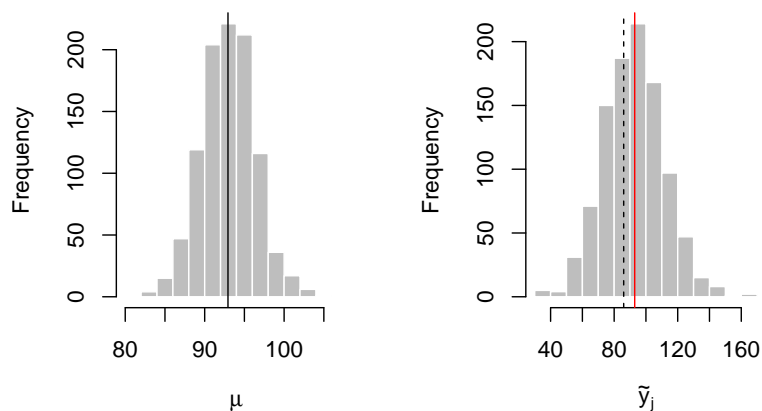


`mcmc_dens_overlay(sims)`



Due to pooling,  $\mathbb{E}[y_6] = \mathbb{E}[y_7] = \mu$  and  $\tilde{y}_j$  has the same distribution independent of  $j$ .

```
# as before, let's just use the first chain
mu <- sims[, 1, 1]; tau <- sims[, 1, 2]
ypred6 <- rnorm(nsamples, mu, tau)
op <- par(mfrow = c(1, 2))
plot_hist(mu, xlab = expression(mu)); abline(v = ybar)
plot_hist(ypred6, xlab = expression(tilde(y)[j]))
abline(v = ybar, col = "red"); abline(v = yj[6], lty = 2)
```



```
par(op)
```

The Stan model is shown below. As in the separate case, the results are equivalent.

```
sm <- stan_model(model_code = "data { int<lower=0> n; vector[n] y; }
parameters { real mu; real<lower=0> tau; }
model { y ~ normal(mu, tau); target += -2 * log(tau); }")
sf <- sampling(sm, data = c("n", "y"))
```

## Hierarchical

This is the model from Section 11.6 we discussed in class:  $y_{ij} | \theta_j, \sigma^2 \stackrel{\text{iid}}{\sim} N(\theta_j, \sigma^2)$ ,  $\theta_j | \mu, \tau^2 \stackrel{\text{iid}}{\sim} N(\mu, \tau^2)$ , and non-informative priors  $\mathbb{P}(\mu, \tau, \log \sigma) \propto 1$ .

```
sims <- mcmc_array(nsamples, nchains,
                  c(paste0("theta", 1:J), "mu", "tau", "sigma", "lp_"))

theta <- yj; mu <- mean(theta)
nt <- S + sum(nj * (yj - theta) ^ 2)
for (chain in 1:nchains) {
  for (it in 1:iter) {
    sigma2 <- rinvchisq(1, n, nt)
    tau2 <- rinvchisq(1, J - 1, sum((theta - mu) ^ 2))
    theta <- rnorm(J, (nj * yj / sigma2 + mu / tau2) / (nj / sigma2 + 1 / tau2),
                  1 / sqrt(nj / sigma2 + 1 / tau2))
    mu <- rnorm(1, mean(theta), sqrt(tau2 / J))
    nt <- S + sum(nj * (yj - theta) ^ 2)
    target <- -n / 2 * log(sigma2) - nt / (2 * sigma2) +
              sum(dnorm(theta, mu, sqrt(tau2), log = TRUE)) - log(sigma2) - .5 * log(tau2)

    if (it > warmup)
      sims[it - warmup, chain, ] <- c(theta, mu, sqrt(tau2), sqrt(sigma2), target)
  }
}
```

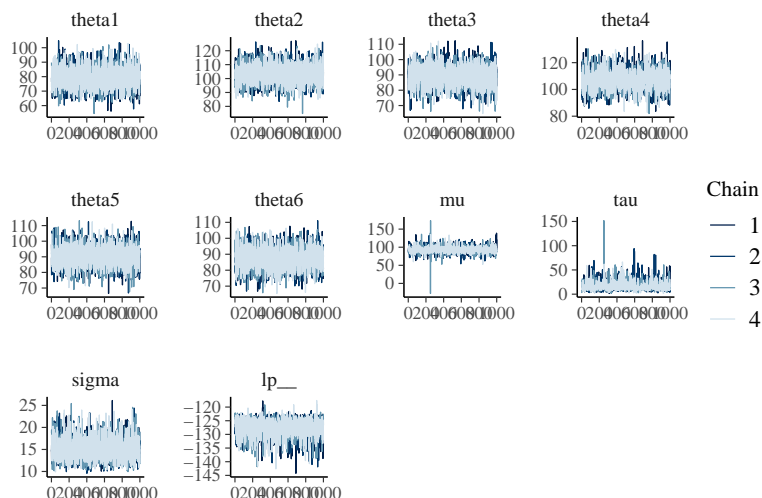
```
monitor(sims, warmup = 0)
```

Inference for the input samples (4 chains: each with iter = 1000; warmup = 0):

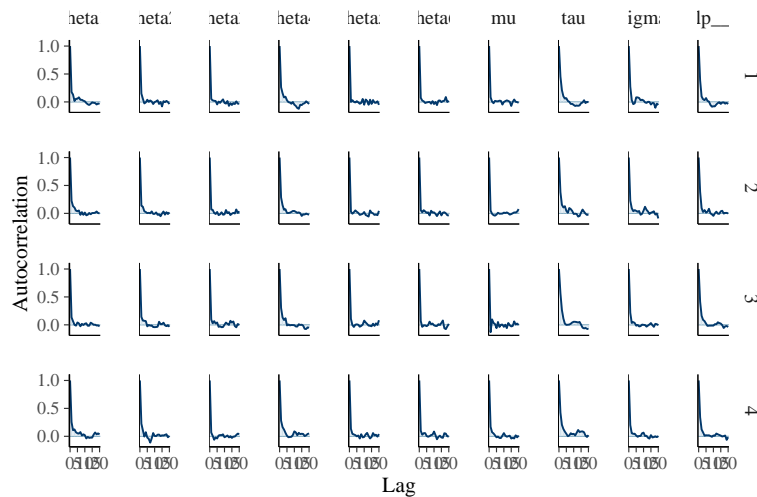
	Q5	Q50	Q95	Mean	SD	Rhat	Bulk_ESS	Tail_ESS
theta1	69.1	79.7	91.4	79.8	6.8	1	1834	2168
theta2	92.9	103.3	113.8	103.3	6.3	1	2480	3082
theta3	78.9	89.1	98.7	89.0	6.1	1	3278	3644
theta4	95.6	107.6	118.7	107.5	7.0	1	1820	1674
theta5	80.6	90.7	100.0	90.6	6.0	1	3371	3776
theta6	77.4	87.5	97.6	87.5	6.1	1	3178	3519
mu	81.3	93.1	105.4	93.1	8.1	1	3072	3136
tau	6.1	14.1	33.0	16.2	9.5	1	1106	911
sigma	11.7	14.7	19.2	14.9	2.3	1	2169	3239
lp__	-133.5	-127.2	-123.5	-127.7	3.1	1	1721	2241

For each parameter, Bulk\_ESS and Tail\_ESS are crude measures of effective sample size for bulk and tail quantities respectively (an ESS > 100 per chain is considered good), and Rhat is the potential scale reduction factor on rank normalized split chains (at convergence, Rhat <= 1.05).

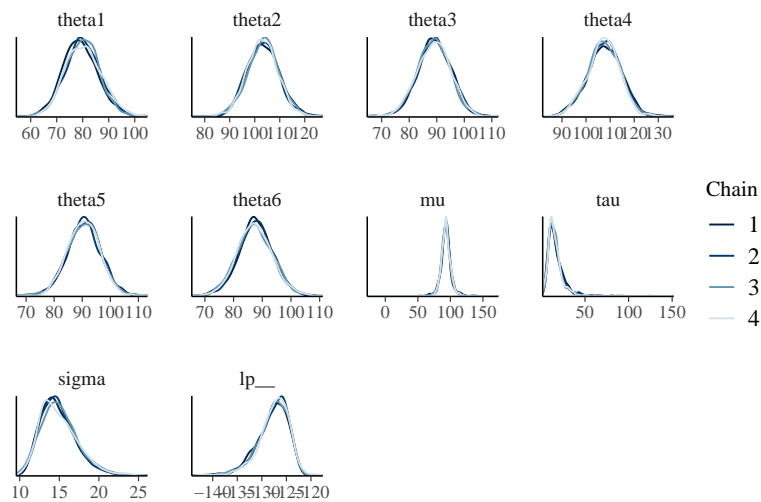
```
mcmc_trace(sims)
```



```
mcmc_acf(sims)
```

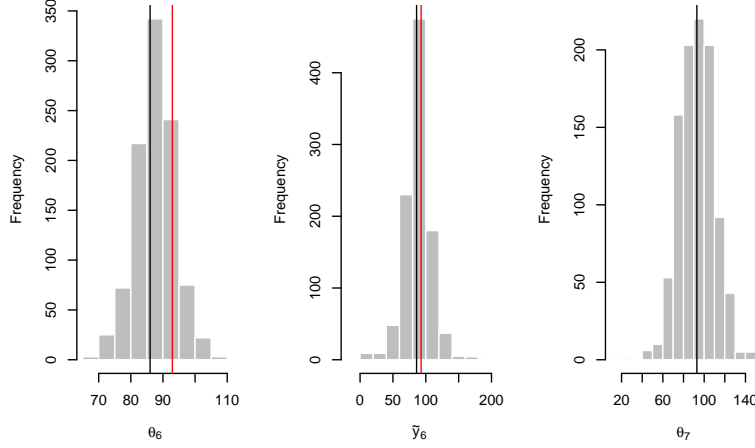


```
mcmc_dens_overlay(sims)
```



The posterior on  $\theta_6$  can be estimated directly from the MCMC samples. For  $\tilde{y}_6$  and  $\theta_7$  we have to use posterior predictive estimates.

```
theta6 <- sims[, 1, 6]; mu <- sims[, 1, 7]
sigma <- sims[, 1, 8]; tau <- sims[, 1, 9]
ypred6 <- rnorm(nsamples, theta6, sigma)
theta7 <- rnorm(nsamples, mu, tau)
op <- par(mfrow = c(1, 3))
plot_hist(theta6, xlab = expression(theta[6]))
abline(v = yj[6]); abline(v = ybar, col = "red")
plot_hist(ypred6, xlab = expression(tilde(y)[6]))
abline(v = yj[6]); abline(v = ybar, col = "red")
plot_hist(theta7, xlab = expression(theta[7]))
abline(v = ybar)
```



`par(op)`

### 3

```
y <- c(13, 52, 6, 40, 10, 7, 66, 10, 10, 14, 16, 4,
       65, 5, 11, 10, 15, 5, 76, 56, 88, 24, 51, 4,
       40, 8, 18, 5, 16, 50, 40, 1, 36, 5, 10, 91,
       18, 1, 18, 6, 1, 23, 15, 18, 12, 12, 17, 3)
n <- length(y)

eps <- 1e-6
boxcox <- function(y, lambda)
  if (abs(lambda) < eps) log(y) else (y ^ lambda - 1) / lambda
```

#### (a)

By variable transformation, we know that

$$\mathbb{P}(y_i | \mu, \sigma^2, \lambda) = \left| \frac{dw_i}{dy_i} \right| \mathbb{P}(w_i | \mu, \sigma^2) = y_i^{\lambda-1} N(w_i | \mu, \sigma^2),$$

and so we have the joint posterior from likelihood independence and by incorporating the prior,

$$\mathbb{P}(\mu, \sigma^2, \lambda | y) \propto \mathbb{P}(y | \mu, \sigma) \mathbb{P}(\mu, \sigma^2, \lambda) = \frac{1}{\sigma} \prod_{i=1}^n N(w_i | \mu, \sigma^2) y_i^{\lambda-1}.$$

#### (b)

Conditional on  $\lambda$  we have the Box-Cox transformed responses  $w_i$ , and so the conditional posterior for  $\mu$  and  $\sigma^2$  depend only on  $w$ . It is straightforward to see that from  $\bar{w} | \mu, \sigma^2 \sim N(\mu, \sigma^2/n)$  and a flat prior on  $\mu$  we have  $\mu | \lambda, \sigma^2, y \sim N(\bar{w}, \sigma^2/n)$  and, from the joint posterior above,

$$\mathbb{P}(\sigma^2 | \mu, \lambda, y) \propto (\sigma^2)^{-((n-1)/2+1)} \exp \left\{ -\frac{\sum_i (w_i(\lambda) - \mu)^2}{2\sigma^2} \right\},$$

that is,  $\sigma^2 | \mu, \lambda, y \sim \text{Inv-}\chi^2(n-1, \sum_i (w_i(\lambda) - \mu)^2/n)$ . For the conditional posterior  $\lambda | \mu, \sigma^2, y$  we have to resort to a random walk MH step. By hand tuning we set the random walk standard deviation at 0.1 and thin out every 10th MCMC sample to reduce autocorrelation.

```

sigma_lambda <- .1
nthin <- 10
sims <- mcmc_array(nsamples, nchains, c("mu", "sigma", "lambda", "lp_"))

for (chain in 1:nchains) {
  lambda <- 0; sigma <- 0
  w <- boxcox(y, lambda)
  for (it in 1:iter) {
    for (thin in 1:nthin) {
      mu <- rnorm(1, mean(w), sigma / sqrt(n))
      sigma <- sqrt(rinvchisq(1, n, sum((w - mu) ^ 2)))
      target <- sum(dnorm(w, mu, sigma, log = TRUE)) +
        (lambda - 1) * sum(log(y))
      # lambda / mu, sigma2:
      lambda_c <- rnorm(1, lambda, sigma_lambda)
      w_c <- boxcox(y, lambda_c)
      target_c <- sum(dnorm(w_c, mu, sigma, log = TRUE)) +
        (lambda_c - 1) * sum(log(y))
      log_r <- target_c - target
      if (log_r >= 0 || log(runif(1)) <= log_r) { # accept?
        lambda <- lambda_c; w <- w_c; target <- target_c
      }
    }

    if (it > warmup)
      sims[it - warmup, chain, ] <- c(mu, sigma, lambda, target - log(sigma))
  }
}

```

(c)

The chains seem to have converged, but with high autocorrelation (even after thinning). Since the 95% credible interval for  $\lambda$  includes 0 but not 0.5, a log transformation is more likely than a square root transformation.

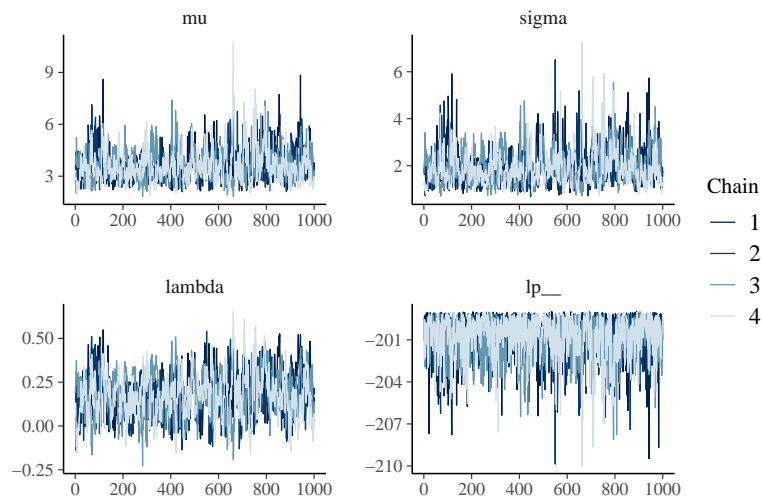
```
monitor(sims, warmup = 0)
```

Inference for the input samples (4 chains: each with iter = 1000; warmup = 0):

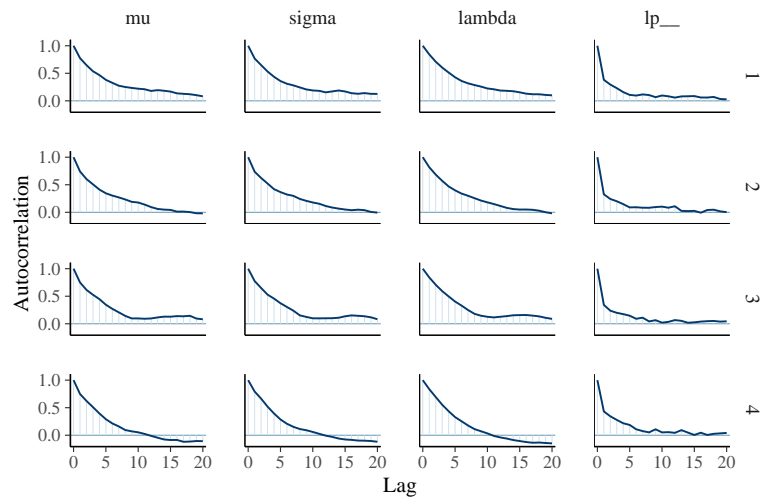
	Q5	Q50	Q95	Mean	SD	Rhat	Bulk_ESS	Tail_ESS
mu	2.5	3.5	5.3	3.6	0.9	1.01	431	795
sigma	1.0	1.8	3.2	1.9	0.7	1.01	417	843
lambda	0.0	0.2	0.4	0.2	0.1	1.01	390	659
lp_	-203.6	-200.4	-199.2	-200.8	1.5	1.00	865	1337

For each parameter, Bulk\_ESS and Tail\_ESS are crude measures of effective sample size for bulk and tail quantities respectively (an ESS > 100 per chain is considered good), and Rhat is the potential scale reduction factor on rank normalized split chains (at convergence, Rhat <= 1.05).

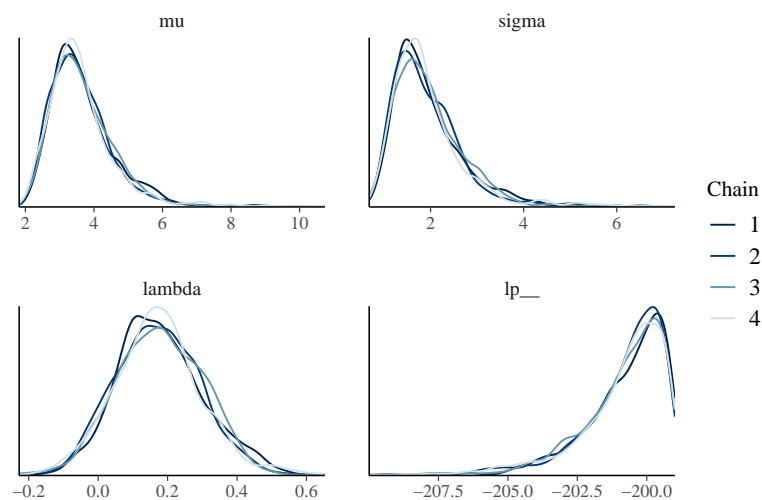
```
mcmc_trace(sims)
```



```
mcmc_acf(sims)
```



```
mcmc_dens_overlay(sims)
```



```
Q <- function(alpha) function(x) quantile(x, alpha)
alpha <- .05
as_tibble(sims[, 1, 1:3]) %>%
```



```
summarize_all(list(lower = Q(alpha / 2), upper = Q(1 - alpha / 2)))
```

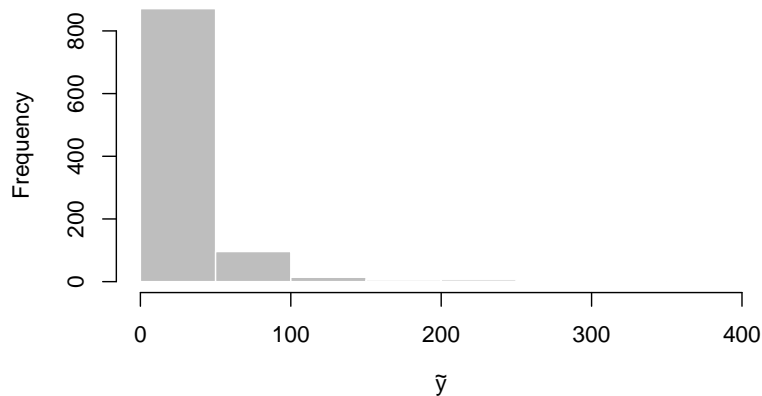
```
# A tibble: 1 x 6
```

	mu_lower	sigma_lower	lambda_lower	mu_upper	sigma_upper	lambda_upper
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	2.39	0.996	-0.0463	5.80	3.85	0.447

(d)

The posterior predictive distribution is smoother and unimodal given the simple model we have adopted, but it seems to be in good agreement, overall, with the data.

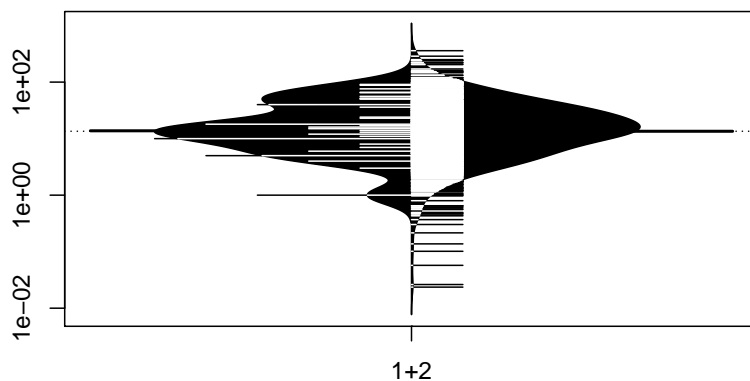
```
mu <- sims[, 1, 1]; sigma <- sims[, 1, 2]; lambda <- sims[, 1, 3]
wpred <- rnorm(nsamples, mu, sigma)
ypred <- (lambda * wpred + 1) ^ (1 / lambda)
plot_hist(ypred, xlab = expression(tilde(y)))
```



```
quantile(ypred, c(.025, .975), na.rm = TRUE)
```

	2.5%	97.5%
	1.126085	105.503216

```
beanplot(y, ypred, side = "both")
```



(extra: Stan)

Here's Stan code for the same model. The results are very similar, but Stan has slightly less autocorrelation.

```
functions {
  real boxcox (real y, real lambda) {
    real eps = 1e-6;
```

```

        if (fabs(lambda) < eps) return log(y);
        return (y ^ lambda - 1) / lambda;
    }
}
data {
    int<lower=0> n;
    vector[n] y;
}
parameters {
    real mu;
    real<lower=0> sigma;
    real lambda;
}
transformed parameters {
    vector[n] w;
    for (i in 1:n) w[i] = boxcox(y[i], lambda);
}
model {
    w ~ normal(mu, sigma);
    target += (lambda - 1) * sum(log(y)); // log Jacobian
}

```